

CS 6480: Lab Assignment #2 (phase 2-3)

Detecting cloud coresidency using hard-disk based covert channel

Binh Nguyen¹

¹School of Computing, University of Utah

November 13, 2013

1 Introduction

As cloud providers allow multi-tenancy in their cloud, hostile virtual machines (VMs) can be placed on the same physical machine that hosts other regular users. The consequence of this is the possibility of cross-VM attacks in which the adversary penetrates the hypervisor or uses side-channel to commit harms to other VMs. This type of attacks often comes with two step: placement and extraction. During placement, the attacker exploits the spatiality and/or locality characteristic of cloud providers' resource allocation mechanism to place its VM on the same physical machine with the victim with a lowest cost. Recent work shows that adversary can achieve up to 40% of co-residency in Amazon EC2 cloud with just a few dollars [1].

Many side-channel attacks exploit the imperfect resource isolation in virtualization. VMs on a same physical machine share local disk, CPU data's cache. By listening to the performance of these shared channels, VMs can eavesdrop each other. This type of eavesdropping was shown be able to extract RSA and AES secret keys, or listen to keystroke from the other VM. To commit these attacks, however, the attackers have to determine whether they have the co-residency with the victim.

This lab assignment leverages side-channel monitoring to develop a hard disk based co-residency

detection. By building a coarse-grained hard disk based covert channel, a VM can listen to another VM's disk actions and determine whether they are co-resident (phase A). The lab assignment takes a step further by demonstrating that an attacker can detect co-residency without having any control over the victim but using a remote machine to cause I/O loads to the victim and listening to the covert channel (phase B).

2 Design

The experiment setup consists of two VMs hosted on one physical machine and sharing the same local disk (as in Figure 1). The idea behind disk based covert channel is that if the two VMs simultaneously read/write the local disk, both should see a degradation on disk performance; if A idles and B reads/writes the disk, B should see a better disk performance compared to the first case. The two distinguishable disk performances is sufficient to present low/high (0/1) signals and can be used for communication between A and B.

When either sender or receiver reads the local disk, it will perceive a high disk reading speed (say N MB/s). When both read the disk in the mean time, since the disk is shared, each reader will perceive a relatively lower disk reading speed

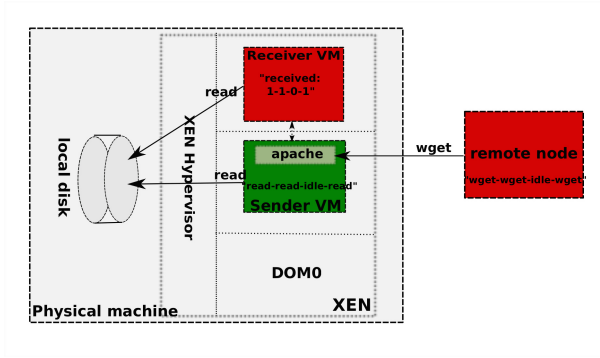


Figure 1: Experiment architecture. Sender and receiver (VMs) reside on the same physical machine that runs XEN. Both VMs share the same local disk. The sender (victim) also runs apache server and allows requests from a remote node.

(theoretically $N/2$ MB/s). If the receiver keeps reading the disk and the sender: (i) reads the disk when it wants to send "1", (ii) idle when it wants to send "0", the receiver will see low disk reading speed when the sender sends "1" and high disk reading speed when the sender sends "0". By keeping the receiver reading the disk and the sender reads or idles, the sender can send "1" or "0" signal to the receiver. This is the disk based covert channel.

The read actions of sender and receiver, however, have to be synchronized for accuracy. While the receiver keeps reading the disk, the sender has to maintain the timing for idle/reading duration so that the receiver can distinguish between low and high period. A possible solution for this timing is to say the sender idle or read per X seconds, and the bit rate of the covert channel is X bps.

For phase B, the idea is to let a remote node to trigger I/O loads on the sender and let the receiver detect that. In this experiment, the remote node issues wget commands to download a large file on the sender. To serve the remote host, the sender has to read local disk for the large file and this causes I/O loads. In a real attack scenario, the remote host trigger the sender at the behest of the attacker, and the trigger-

ing pattern can be used to detect co-residency. For example, if the remote host periodically gets the file and idles, the receiver (if co-resides with the sender) should see "1010..." signal.

3 Implementation

This experiment runs on Emulab [2] machines. One physical node, which is able to do virtualization, hosts two XEN VMs instances. The remote node is hosted on a separate machine which is able to connect to the other machine.

To measure disk reading speed, the experiment uses hdparm, a linux built-in command that continuously reads a disk for about three seconds and reports the total amount of data read as well as the reading speed. The receiver periodically captures disk performance using hdparm commands (every 10 seconds). The sender receives a "signal" string and sends it using idling or reading actions (also every 10 seconds). For example, the sender is commanded to send "11010", for every 10 seconds, it performs: read (1), read (1), idle (0), read (1), idle (0). The receiver sees "11010" correspondingly.

For phase B, the sender runs a Apache server and allows the remote host to download a large file stored in the local disk. The remote host trigger the download to send a "1" signal to the receiver and idle to send "0".

4 Result

The experiment successfully implemented a hard disk based covert channel with a bit rate of 6 bps with the accuracy of nearly 100%. The sender is able to send binary signal to the receiver (Figure 2 and 4).

The actual receiver's disk reading speed and received signal are shown in figure ???. The threshold to determine 1 or 0 signal is shown in blue line. This threshold determined as 0.85 of the disk speed mea-

```

binh6480@vm0:~/6480/lab2/phase2_3$ ./sender 1101101
Sender writing 1101101 ...
SENDER 00:57:40 : writing 1 ...
SENDER 00:57:50 : writing 1 ...
SENDER 00:58:00 : writing 0 ...
SENDER 00:58:10 : writing 1 ...
SENDER 00:58:20 : writing 1 ...
SENDER 00:58:30 : writing 0 ...
SENDER 00:58:40 : writing 1 ...
SENDER 00:58:46 : finished writing ['1', '1', '0', '1', '1', '0', '1']. Exit
Finished writing, exit.
binh6480@vm0:~/6480/lab2/phase2_3$

```

Figure 2: Sender's screen shoot. Sender sends "1101101" signal, 1 bit every 10 seconds.

```

binh6480@vm1:~/6480/lab2/phase2_3$ ./receiver
Receiver is training to determine LOW/HIGH values ...
LOW/HIGH threshold = 57.596 MB/s
Started listening:
=====
00:57:46 RECEIVER : speed= 30.43 MB/s , OUTPUT = 1.
00:57:56 RECEIVER : speed= 31.11 MB/s , OUTPUT = 1.
00:58:06 RECEIVER : speed= 68.82 MB/s , OUTPUT = 0.
00:58:16 RECEIVER : speed= 29.34 MB/s , OUTPUT = 1.
00:58:26 RECEIVER : speed= 27.40 MB/s , OUTPUT = 1.
00:58:36 RECEIVER : speed= 69.38 MB/s , OUTPUT = 0.
00:58:46 RECEIVER : speed= 27.08 MB/s , OUTPUT = 1.
00:58:56 RECEIVER : speed= 68.70 MB/s , OUTPUT = 0.

```

Figure 3: Receiver's screen shoot. Receiver receives "1101101" signal, 1 bit every 10 seconds.

sured at the time the receiver starts.

In phase B, the remote host trigger "1010" signal and that signal was picked up by the receiver (Figure 5 and 6). In phase B, the remote host downloads a file (or idles) in 60-seconds periods, and the receiver reads the disk at the same frequency.

5 Discussion

Although the experiment successfully implemented a disk based covert channel, the channel has some limitations with speed and accuracy. To increase speed of the channel, the measurements have to happen more frequently. The result of this is less accurate measurements since hard disk is not responsive enough. Also, sender and receiver have to be synchronized for better accuracy. Using built-in command (hdparm), the experiment was

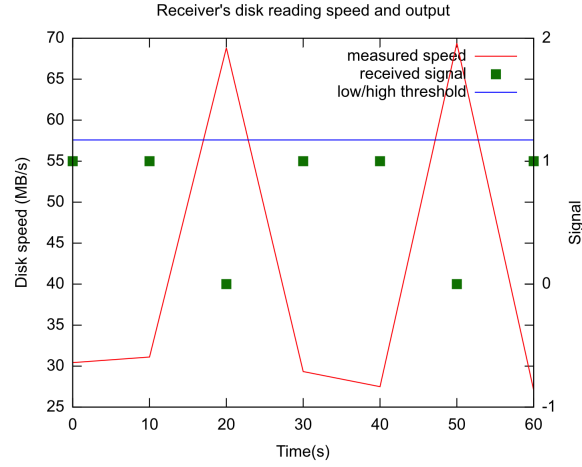


Figure 4: Receiver's disk reading speed and received signal

```

binh6480@othernode:~/6480/lab2/phase2_3$ ./trigger
**Trigger: 1010 ...
Synchronizing. Please wait...
TRIGGER: SENDING 1 ...
TRIGGER: SENDING 0 ...
TRIGGER: SENDING 1 ...
TRIGGER: SENDING 0 ...
DONE!
binh6480@othernode:~/6480/lab2/phase2_3$

```

Figure 5: Remote host's screen shoot. Remote host triggers "1010" signal which corresponding to "download-idle-download-idle" series of actions.

```

binh6480@vm1:~/6480/lab2/phase2_3$ ./receiver.p3
RECEIVER: Preparing, please wait ...
LOW/HIGH threshold = 58.1031666667 MB/s
Started listening:
=====
22:28:00 RECEIVER : OUTPUT = 1.
22:29:00 RECEIVER : OUTPUT = 0.
22:30:00 RECEIVER : OUTPUT = 1.
22:31:00 RECEIVER : OUTPUT = 0.
22:32:00 RECEIVER : OUTPUT = 0.

```

Figure 6: Receiver's screen shoot in phase B. As the remote host downloads a large file from sender, the receiver perceives disk loads and detects "1010", the same pattern as the remote host.

able to implement a highly accurate covert channel. However, since `hdparm` tends to stress test the disk for a sufficient amount of time (3s), the covert channel could not be faster than 20 bps.

The implementation of phase B is not guaranteed 100% successful. The magnitude of the I/O loads that the remote host causes on the sender depends on (i) the download link's speed and (ii) the local disk speed. If the download link's speed is too low and/or the local disk speed is too high, then downloading might not be significant enough to cause disk performance degradation on the local disk. In the experiment, if local disk speed is about 100 MB/s without load, and the download speed is about 10 MB/s, the degradation of disk performance is less than 10% (reading speed is always higher than 90 MB/s regardless of the downloading).

The experiment is done using Python and Bash scripts. All the scripts for this lab assignment could be found here.

References

- [1] RISTENPART, T., TROMER, E., SHACHAM, H., AND SAVAGE, S. Hey, you, get off of my cloud: exploring information leakage in third-party compute clouds. In *Proceedings of the 16th ACM conference on Computer and communications security* (2009), ACM, pp. 199–212.
- [2] WHITE, B., LEPREAU, J., STOLLER, L., RICCI, R., GURUPRASAD, S., NEWBOLD, M., HIBLER, M., BARB, C., AND JOGLEKAR, A. An integrated experimental environment for distributed systems and networks. *ACM SIGOPS Operating Systems Review* 36, SI (2002), 255–270.